

CS5340: Uncertainty Modeling in AI

Coding Assignment 2

Deadline: April 19th, 2019

Problem 1. (Image Recovery using Hopfield Network)

Hopfield Networks

A Hopfield network (HopNet) is a fully-connected Ising model with a symmetric weight matrix, i.e., the weight matrix has the following properties:

- $\mathbf{W}_{ii} = 0$
- $\mathbf{W}_{ij} = \mathbf{W}_{ji}$




Figure 1: A Hopfield Network with 4 nodes.

A HopNet also has a bias vector b . The parameters ($\theta = [\mathbf{W}, b]$) can be learned from data. HopNets can be used as an *associative memory*. We train the HopNet on a fully observed set corresponding to some patterns that we want to memorize. At test time we present a partial or corrupted pattern and the network attempts to complete the pattern. Once the parameters are learned, pattern completion can be performed using iterated conditional modes or ICM (similar to image denoising example).

The conditional probability of a node taking the value 1 is given by:

$$p(x_i = 1 | \mathbf{x}_{-i}, \theta) = \sigma(\mathbf{W}_{i,:}^T \mathbf{x}_{-i} + b_i)$$

where σ is the sigmoid function, i.e., $\sigma(x) = \frac{1}{1 + \exp(-x)}$.

For a more detailed description, refer “Machine Learning – A Probabilistic Perspective”, Kevin Murphy (Chapter 19, 19.4.2) and the Wikipedia article: Hopfield Network.

Goal

You are expected to implement a Hopfield network for image recovery. Given a set of binary images, your goal is to:

- Learn the parameters θ using a set of training images.
- Predict the most likely assignment of pixel values for the corrupted versions of training images using ICM.



Figure 2: Example result.

Parameter Learning (15 Marks)

Your task is to learn the parameters using two methods:

1. **Hebbian Learning Rule (5 Marks):** This rule is often stated as, “Neurons that fire together, wire together. Neurons that fire out of sync, fail to link”. Complete the function `learn_hebbian(imgs)` in the skeleton code. The function takes a numpy array of shape $(n, 32, 32)$ where n is the number of 32×32 binary images. The function should return a tuple (W, b) where W is the learned 1024×1024 weight matrix and b is the 1024 dimensional bias vector.

Allowed libraries: `numpy` and `scipy`.

2. **Maximum Pseudo-likelihood (10 Marks):** The parameters can also be learned by using gradient-based methods to maximize the *pseudo-likelihood* which is given by

$$PL(\theta) = \prod_{k=1}^N \prod_{i=1}^D p(x_i^k | \mathbf{x}_{-i}^k, \theta)$$

Note that this is not the actual likelihood which is given by:

$$L(\theta) = \prod_{k=1}^N p(\mathbf{x}^k | \theta)$$

Complete the function `learn_maxpl(imgs)` in the skeleton code. The function takes a numpy array of shape $(n, 32, 32)$ as the input, where n is the number of 32×32 binary images. The function should return a tuple (W, b) , where W is the learned 1024×1024 weight matrix and b is the 1024 dimensional bias vector.

Allowed libraries: `numpy` and `scipy`. Libraries `torch`, `autograd`, etc., are also allowed for automatic gradient computation, if required.

Image Recovery (5 Marks)

Complete the function `recover(cimgs, W, b)` in the skeleton code. The function takes a numpy array of shape $(n, 32, 32)$ and W and b as the input, where n is the number of 32×32 corrupted binary images, and W and b are the learned weight matrix and bias vector respectively. The function should return a numpy array of shape $(n, 32, 32)$ which contains the images recovered using the HopNet.

Submission Format

Submit only the python (.py) file renamed to `YourMatricNumber-PartnerMatricNumber.py` on IVLE. If your matric number is A0174067B and your partner's is A0175067A, then the file should be named `A0174067B-A0175067A.py`. If you're doing the assignment as an individual, name it as `YourMatricNumber.py`. Submit **only one** python file per group.

Code Skeleton

You are only allowed to modify the functions `learn_hebbian(imgs)`, `learn_maxpl(imgs)`, and `recover(cimgs, W, b)` in the code skeleton. Adding helper functions is allowed but modifying `main()` in any way is not allowed.

```
1  '''
2  Description: CS5340 - Hopfield Network
3  Name: Your Name, Your partner's name
4  Matric No.: Your matric number, Your partner's matric number
5  '''
6
7
8  import matplotlib
9  matplotlib.use('Agg')
10 import numpy as np
11 import glob
12 import matplotlib.pyplot as plt
```

```

13 from PIL import Image, ImageOps
14
15
16 def load_image(fname):
17     img = Image.open(fname).resize((32, 32))
18     img_gray = img.convert('L')
19     img_eq = ImageOps.autocontrast(img_gray)
20     img_eq = np.array(img_eq.getdata()).reshape((img_eq.size[1], -1))
21     return img_eq
22
23
24 def binarize_image(img_eq):
25     img_bin = np.copy(img_eq)
26     img_bin[img_bin < 128] = -1
27     img_bin[img_bin >= 128] = 1
28     return img_bin
29
30
31 def add_corruption(img):
32     img = img.reshape((32, 32))
33     t = np.random.choice(3)
34     if t == 0:
35         i = np.random.randint(32)
36         img[i:(i + 8)] = -1
37     elif t == 1:
38         i = np.random.randint(32)
39         img[:, i:(i + 8)] = -1
40     else:
41         mask = np.sum([np.diag(-np.ones(32 - np.abs(i)), i)
42                        for i in np.arange(-4, 5)], 0).astype(np.int)
43         img[mask == -1] = -1
44     return img.ravel()
45
46
47 def learn_hebbian(imgs):
48     img_size = np.prod(imgs[0].shape)
49     #####
50     #####
51     weights = np.zeros((img_size, img_size))
52     bias = np.zeros(img_size)
53     # Complete this function
54     # You are allowed to modify anything between these lines
55     # Helper functions are allowed
56     #####
57     #####

```

```

58     return weights, bias
59
60
61 def learn_maxpl(imgs):
62     img_size = np.prod(imgs[0].shape)
63     #####
64     #####
65     weights = np.zeros((img_size, img_size))
66     bias = np.zeros(img_size)
67     # Complete this function
68     # You are allowed to modify anything between these lines
69     # Helper functions are allowed
70     #####
71     #####
72     return weights, bias
73
74
75 def plot_results(imgs, cimgs, rings, fname='result.png'):
76     '''
77     This helper function can be used to visualize results.
78     '''
79     img_dim = 32
80     assert imgs.shape[0] == cimgs.shape[0] == rings.shape[0]
81     n_imgs = imgs.shape[0]
82     fig, axn = plt.subplots(n_imgs, 3, figsize=[8, 8])
83     for j in range(n_imgs):
84         axn[j][0].axis('off')
85         axn[j][0].imshow(imgs[j].reshape(img_dim, img_dim), cmap='Greys_r')
86         axn[0, 0].set_title('True')
87     for j in range(n_imgs):
88         axn[j][1].axis('off')
89         axn[j][1].imshow(cimgs[j].reshape(img_dim, img_dim), cmap='Greys_r')
90         axn[0, 1].set_title('Corrupted')
91     for j in range(n_imgs):
92         axn[j][2].axis('off')
93         axn[j][2].imshow(rings[j].reshape((img_dim, img_dim)), cmap='Greys_r')
94         axn[0, 2].set_title('Recovered')
95     fig.tight_layout()
96     plt.savefig(fname)
97
98
99 def recover(cimgs, W, b):
100     img_size = np.prod(cimgs[0].shape)
101     #####
102     #####

```

```

103     rings = []
104     # Complete this function
105     # You are allowed to modify anything between these lines
106     # Helper functions are allowed
107     #####
108     #####
109     return rings
110
111
112 def main():
113     # Load Images and Binarize
114     ifiles = sorted(glob.glob('images/*'))
115     timgs = [load_image(ifile) for ifile in ifiles]
116     imgs = np.asarray([binarize_image(img) for img in timgs])
117
118     # Add corruption
119     cimgs = []
120     for i, img in enumerate(imgs):
121         cimgs.append(add_corruption(np.copy(imgs[i])))
122     cimgs = np.asarray(cimgs)
123
124     # Recover 1 -- Hebbian
125     Wh, bh = learn_hebbian(imgs)
126     rings_h = recover(cimgs, Wh, bh)
127     np.save('hebbian.npy', rings_h)
128
129     # Recover 2 -- Max Pseudo Likelihood
130     Wmpl, bmpl = learn_maxpl(imgs)
131     rings_mpl = recover(cimgs, Wmpl, bmpl)
132     np.save('mpl.npy', rings_mpl)
133
134
135 if __name__ == '__main__':
136     main()

```
